Formalizing Cohomology for Multi-Agent Systems: A Categorical Approach to Economic and Reinforcement Learning Models

Co-advisors: Prof. Fernando Tohmé, Prof. Neil Ghani (Strathclyde) and Dr. Toby St Clere Smithe (Topos Institute)
Committee Members: Dr. Hans Riess (Georgia Tech), Prof. Mike Shulman (University of San Diego), Prof. Colin McLarty (Case Western), Prof. Rocco Gangle (Endicott) and Prof. William Mansky (UIC)

Eric Schmid

April 7, 2025

Abstract

The explosion of complex multi-agent systems across disciplines demands new mathematical tools to understand their emergent properties. This research develops a framework rooted in sheaf theory and category-theoretic foundations to tackle this challenge head-on. I propose to create rigorous formalizations of sheaves in dependently typed languages and apply them directly to economic and reinforcement learning models. By establishing a novel "isomorphism" between multi-agent systems and sheaf cohomology, my work aims to build a powerful framework for analyzing complex systems at multiple scales. This approach bridges the often-disconnected worlds of theoretical computer science, economic modeling, and machine learning, with implementations allowing practical applications in these domains. The ultimate goal is both theoretical advancement and computational tools that address real-world complex systems challenges.

1 Introduction and Background

1.1 Motivation

Whether analyzing financial markets, social networks, or multi-agent AI systems, our tools for capturing interactions between agents, information flow, and emergent properties remain frustratingly limited. The gap is particularly evident when trying to maintain computational efficiency while preserving theoretical rigor. My research tackles this problem by turning to sheaf theory – a branch of mathematics explicitly designed to relate local and global properties.

The insight driving this research comes from a surprising observation: many problems in economics and reinforcement learning share a fundamental characteristic. Local decisions made by individual agents with limited information collectively determine global system behavior. This mirrors precisely the mathematical domain where sheaf theory excels – analyzing how local data and constraints generate global structures. While economists, computer scientists, and mathematicians have all circled these problems, they've largely done so independently, missing opportunities for cross-pollination.

Sheaf theory has already demonstrated its power in mathematics and physics, but its application to computational contexts remains surprisingly underdeveloped. The time is ripe to bridge this gap, particularly for multi-agent and economic systems where its local-to-global perspective offers a natural fit. I'm proposing to develop a formalized sheaf-theoretic approach to modeling complex systems of interacting agents, with explicit applications to economic markets and reinforcement learning frameworks.

1.2 Related Work

My research builds on several pioneering efforts spanning mathematics, computer science, and economics:

- Sterling's Agda implementation of sheaves has broken ground in constructive sheaf semantics, but hasn't been extended to multi-agent systems
- Escardo's type-topology library in univalent foundations provides essential tools for homotopy type theory that can inform our approach
- Recent work in compositional game theory has revealed the potential of category-theoretic methods in economics, though without leveraging sheaf cohomology
- Advances in languages like Agda and Idris have made formal verification of complex mathematical structures increasingly feasible

My preliminary research suggests a striking gap – despite sheaf cohomology's natural fit for multi-agent systems, direct applications to economics or reinforcement learning remain scarce. This presents both a challenge and an opportunity to develop novel frameworks with significant theoretical and practical impact.

2 Research Questions and Objectives

My research confronts four questions:

1. What novel insights emerge when we examine multi-agent systems through the lens of sheaf cohomology? Can we characterize emergent properties that traditional approaches miss?
2. How can we effectively represent economic and reinforcement learning systems as sheaves? What's the optimal mapping between system components and sheaf-theoretic constructs?

3. How can we formalize these sheaf-theoretic representations in dependently typed languages to create verified computational foundations?
4. How do we translate these formalizations into computationally efficient implementations that scale to real-world applications?

To address these questions, I'll pursue five core objectives:

1. Develop a cohomological interpretation of system properties that reveals new insights about emergent behaviors in multi-agent systems
2. Establish formal mappings between multi-agent system components and sheaf-theoretic constructs, with specific instantiations for economic modeling and reinforcement learning
3. Create a comprehensive formalization of the relevant sheaf structures in a dependently typed language
4. Implement efficient computational methods for working with these structures, exploring parallel evaluation using interaction combinators
5. Demonstrate practical applications through concrete case studies in economics and reinforcement learning

3 Methodology and Research Plan

I'll tackle this ambitious agenda through five phases:

3.1 Phase 1: Survey and Analysis of Existing Formalizations (6 months)

I'll begin with a deep dive into:

- Formalizations in category theory, sheaf theory, and topology
- Sheaf neural networks (Hansen and Gebhart)
- Persistent homology and Topological Data Analysis (Carlsson)
- Information theory through group cohomology (Vigneaux)
- Applications of sheaf theory to multi-agent systems
- Implementations in dependently typed languages
- Parallel evaluation architectures like Taelin's HVM2

I'll critically evaluate how these disparate approaches might complement or challenge each other. I'll start by immersing myself in Hansen and Gebhart's work on sheaf neural networks, Carlsson's pioneering contributions to persistent homology, and Vigneaux's information-theoretic models using group cohomology. Beyond surveying existing work, I'll assess implementations like Sterling's sheaf formalization and Escardo's type-topology library to identify leverageable components. I'm particularly interested in exploring how Taelin's interaction combinator approach might revolutionize parallel evaluation of sheaf cohomology computations.

3.2 Phase 2: Identification of Target Applications (6 months)

Armed with theoretical foundations from Phase 1, I'll identify high-impact applications:

- Map economic modeling and reinforcement learning challenges that could benefit from sheaf-theoretic approaches
- Identify problems where local-to-global properties are crucial (market interactions, information aggregation, etc.)
- Develop preliminary sheaf-theoretic formulations
- Establish evaluation criteria for successful formulations

This phase will yield concrete target problems in economics and reinforcement learning that will drive my subsequent formalization work. I'll focus on applications where sheaf representations offer genuine insights or computational advantages that traditional approaches miss.

3.3 Phase 3: Sheaf-Theoretic Modeling Framework (12 months)

Building on my work in Phases 1 and 2, I'll develop a comprehensive framework for modeling multi-agent systems using sheaf theory:

- Map system components (agents, information, decisions) to sheaf-theoretic constructs
- Define mechanisms by which global properties emerge from local behaviors through sheaf cohomology
- Formalize information flows and decision-making as sheaf morphisms
- Develop specific models for the target applications identified earlier

This phase will yield a formal mathematical framework establishing the "isomorphism" between multi-agent systems and sheaf-theoretic constructs, with concrete instantiations for my target applications.

3.4 Phase 4: Formalization in Dependently Typed Languages (6 months)

I'll implement the mathematical framework from Phase 3 in a dependently typed language:

- Use the Curry-Howard isomorphism to encode sheaves and related categorical constructs in Agda or Idris
- Verify critical mathematical properties through formal proofs
- Implement instances of the framework for target applications
- Establish foundations for computational experiments

3.5 Phase 5: Computational Implementation (6 months)

I'll translate theoretical formalization into efficient implementations, exploring several approaches:

1. **Sheaf cohomology implementation**: Algorithms for computing sheaf cohomology over multi-agent systems
2. **Persistent homology / TDA approach**: Topological data analysis methods as alternatives or complements to sheaf cohomology
3. **Information-theoretic implementations**: Approaches based on Vigneaux's group cohomology framework for Shannon entropy
4. **HVM compilation**: Pathways to interaction combinators for efficient parallel evaluation via Taelin's HVM2 architecture
5. **Hybrid approach**: Agda formalization with optimized Python implementations

This phase involves comparative evaluation of these approaches, assessing computational efficiency, theoretical fidelity, and applicability to my target problems. While I'll emphasize the HVM approach for its parallel computation benefits, I'll also explore how persistent homology and information-theoretic implementations might offer unique advantages for specific domains.

3.6 Phase 6: Application and Evaluation (6 months)

I'll apply my framework to concrete problems in:

- Economic modeling, potentially focusing on micro-economy integration or market actor connections
- Reinforcement learning systems, emphasizing incentives, policy formulation, and metrics
- Evaluation of effectiveness, computational efficiency, and theoretical insights

4 Expected Contributions

My research aims to make five significant contributions:

1. Novel theoretical mappings between multi-agent systems and sheaf-theoretic constructs – providing new ways to understand complex systems
2. A verified formalization of sheaves and sheaf cohomology in dependently typed languages – bridging pure mathematics and computational implementation
3. Efficient computational implementations suitable for practical applications – moving beyond theoretical models to usable tools
4. Demonstrated applications in economics and reinforcement learning – showing real-world relevance
5. A foundation for future work on category-theoretic approaches to complex systems – opening new research directions

The core scientific contribution will be establishing an "isomorphism" between specific applications and concepts in sheaf cohomology, enabling rigorous representation of economic and reinforcement learning systems in the language of sheaves.

5 Feasibility and Focus

I recognize the ambitious scope of this research. To make it tractable, I'll:

- Identify the highest-impact subset of the broader conceptual schema with both theoretical significance and practical computational importance
- Leverage existing formalizations wherever possible to avoid reinventing the wheel
- Establish clear evaluation criteria for each phase to maintain focus
- Remain flexible about implementation approaches based on technical feasibility

While the complete integration of sheaf theory into multi-agent system modeling represents my ultimate vision, this PhD will focus on establishing the foundational framework and demonstrating its application in selected high-value contexts. I'll identify promising directions for future exploration beyond the scope of this dissertation.

6 Timeline

- **Year 1:** Complete Phases 1 and 2
  - Months 1-6: Survey and analysis of existing formalizations
  - Months 7-12: Identification of target applications in economics and RL
- **Year 2:** Complete Phase 3 and begin Phase 4
  - Months 1-12: Development of sheaf-theoretic modeling framework
  - Months 10-12: Begin formalization in dependently typed languages
- **Year 3:** Complete Phases 4, 5, and 6, thesis writing
  - Months 1-3: Complete formalization in dependently typed languages
  - Months 4-9: Computational implementation and application
  - Months 7-12: Evaluation and thesis writing

7 Conclusion

This proposal charts a course through the unexplored territory where sheaf theory meets multi-agent systems. By bridging abstract mathematics with practical computational implementations, I aim to create a conceptual framework and technological toolset capable of transforming how we simulate and understand complex systems. If successful, this research will not only advance theoretical understanding but also provide practical tools for modeling, analyzing, and optimizing multi-agent systems across domains – from economic markets to artificial intelligence.

The increasing complexity of our world demands new mathematical frameworks. Sheaf theory, with its emphasis on local-to-global relationships, offers a promising path forward. I believe this research represents an opportunity to make a lasting contribution at the intersection of mathematics, economics, and computer science.

References

[1] Sterling, J. Constructive Sheaf Semantics. `https://github.com/jonsterling/constructive-sheaf-semantics`.

[2] Escardo, M. Type-Topology in Univalent Foundations. `https://www.cs.bham.ac.uk/~mhe/agda-new/`.

[3] The Univalent Foundations Program. Homotopy Type Theory: Univalent Foundations of Mathematics. Institute for Advanced Study, 2013.

[4] Taelin, V. HVM2: A Parallel Evaluator for Interaction Combinators. Work in progress, 2024.

[5] Mac Lane, S., & Moerdijk, I. Sheaves in Geometry and Logic: A First Introduction to Topos Theory. Springer, 1994.

[6] Lafont, Y. Interaction Nets. POPL '90, 1990.

[7] Lafont, Y. Interaction Combinators. Information and Computation, 137(1):69–101, 1997.

[8] Mazza, D. A Denotational Semantics for the Symmetric Interaction Combinators. Mathematical Structures in Computer Science, 17:527–62, 2007.

[9] Hansen, J., & Gebhart, T. Sheaf Neural Networks. arXiv:2012.06333, 2020.

[10] Carlsson, G. Topology and data. Bulletin of the American Mathematical Society, 46(2):255-308, 2009.

[11] Vigneaux, J. P. Topology of Statistical Systems: A Cohomological Approach to Information Theory. Springer, 2019.

[12] Bodnar, C., Di Giovanni, F., Chamberlain, B. P., Lió, P., & Bronstein, M. M. Neural Sheaf Diffusion: A Topological Perspective on Heterophily and Oversmoothing in GNNs. arXiv:2202.04579, 2022.

[13] Curry, J. M. Sheaves, Cosheaves, and Applications. University of Pennsylvania, 2014.